

Technological Institute of the Philippines	Quezon City - Computer Engineering
Course Code:	CPE 313
Code Title:	Advanced Machine Learning and Deep Learning
1st Semester	AY 2024-2025
<u>ACTIVITY NO. 4</u>	Edge and Contour Detection
Name	De los Reyes, Jann Moises Nyll
Section	CPE32S3
Date Performed:	February 21, 2025
Date Submitted:	February 21, 2025
Instructor:	Engr. Roman M. Richard

1. Objectives

This activity aims to introduce students to the use of OpenCV for edge detection and contour detection.

2. Intended Learning Outcomes (ILOs)

After this activity, the students should be able to:

- Explain fundamental idea of convolution and the kernel's importance.
- Use different image manipulation methods by using openCV functions.
- Perform contour and edge line drawing on their own images.

3. Procedures and Outputs



```
In [1]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [2]: import numpy as np
import cv2

# This is the file URL: https://media.wired.com/photos/5cdefb92b86e041493d389df/4:3/w_1330,h_998,c_limit/Cul
src = cv2.imread('/content/drive/MyDrive/CVdata/grumpy-cat.png')
dst = cv2.cvtColor(src, cv2.COLOR_BGR2GRAY)
```

Edges play a major role in both human and computer vision. We, as humans, can easily recognize many object types and their pose just by seeing a backlit silhouette or a rough sketch. Indeed, when art emphasizes edges and poses, it often seems to convey the idea of an archetype, such as Rodin's The Thinker or Joe Shuster's Superman.




Software, too, can reason about edges, poses, and archetypes.

OpenCV provides many edge-finding filters, including `Laplacian()`, `Sobel()`, and `Scharr()`. These filters are supposed to turn non-edge regions to black while turning edge regions to white or saturated colors. However, they are prone to misidentifying noise as edges. This flaw can be mitigated by blurring an image before trying to find its edges. OpenCV also provides many blurring filters, including `blur()` (simple average), `medianBlur()`, and `GaussianBlur()`. The arguments for the edge-finding and blurring filters vary but always include `ksize`, an odd whole number that represents the width and height (in pixels) of a filter's kernel.

For blurring, let's use `medianBlur()`, which is effective in removing digital video noise, especially in color images. For edge-finding, let's use `Laplacian()`, which produces bold edge lines, especially in grayscale images. After applying `medianBlur()`, but before applying `Laplacian()`, we should convert the image from BGR to grayscale.

Once we have the result of `Laplacian()`, we can invert it to get black edges on a white background. Then, we can normalize it (so that its values range from 0 to 1) and multiply it with the source image to darken the edges. Let's implement this approach:

```
In [3]: from google.colab.patches import cv2_imshow
cv2_imshow(src)
```



In [4]: `cv2_imshow(dst)`



```
In [5]: def strokeEdges(src, dst, blurKsize = 7, edgeKsize = 5):
        if blurKsize >= 3:
            blurredSrc = cv2.medianBlur(src, blurKsize)
            graySrc = cv2.cvtColor(blurredSrc, cv2.COLOR_BGR2GRAY)
        else:
            graySrc = cv2.cvtColor(src, cv2.COLOR_BGR2GRAY)

        cv2.Laplacian(graySrc, cv2.CV_8U, graySrc, ksize = edgeKsize)
        normalizedInverseAlpha = (1.0 / 255) * (255 - graySrc)
        channels = cv2.split(src)

        for channel in channels:
            channel[:] = channel * normalizedInverseAlpha

        return cv2.merge(channels, dst)
```


For the function above, provide an analysis:

- Try to run the function and pass values for its parameters.
- Change the values in the kSize variables. What do you notice?


Make sure to add codeblocks underneath this section to ensure that your demonstration of the procedure and answers are easily identifiable

```
In [6]: # Sample


new_img = strokeEdges(src, dst)
cv2.imshow(new_img)
```

```
In [7]: #Sample B  
  
B_img = strokeEdges(src,dst,blurKsize= 2)  
cv2_imshow(B_img)
```


```
In [8]: #sample C  
  
C_img = strokeEdges(src,dst,blurKsize=25)  
cv2_imshow(C_img)
```

```
In [9]: #Sample D  
  
D_img = strokeEdges(src,dst,edgeKsize=15)  
cv2_imshow(D_img)
```




```
In [10]: #Sample E  
  
E_img = strokeEdges(src,dst,edgeKsize=5)  
cv2_imshow(E_img)
```



**Observation:**

Based on the following sample testing of various Ksize values above, we notice that there's a change in each parameter. The blurSize determines the amount of blurring applied in the image while edgeSize controls the size of our edge detection. In blurSize, the higher the value, the edge will appear more smoother since it increases its blurring; decreasing it on the other hand will lead to no more detailed and possibly noisy edges to be detected.

Increasing the edgeSize will make our edges thicker while decreasing it makes our edges become thinner and finer.

Custom Kernels -- Getting Convoluted

As we have just seen, many of OpenCV's predefined filters use a kernel. Remember that a kernel is a set of weights, which determine how each output pixel is calculated from a neighborhood of input pixels. Another term for a kernel is a convolution matrix. It mixes up or convolves the pixels in a region. Similarly, a kernel-based filter may be called a convolution filter.



OpenCV provides a very versatile `filter2D()` function, which applies any kernel or convolution matrix that we specify. To understand how to use this function, let's first learn the format of a convolution matrix. It is a 2D array with an odd number of rows and columns. The central element corresponds to a pixel of interest and the other elements correspond to the neighbors of this pixel. Each element contains an integer or floating point value, which is a weight that gets applied to an input pixel's value.

```
In [11]: # Example
kernel = np.array([[-1, -1, -1],
                  [-1, 9, -1],
                  [-1, -1, -1]])
```

Here, the pixel of interest has a weight of 9 and its immediate neighbors each have a weight of -1. For the pixel of interest, the output color will be nine times its input color minus the input colors of all eight adjacent pixels. If the pixel of interest is already a bit different from its neighbors, this difference becomes intensified. The effect is that the image looks sharper as the contrast between the neighbors is increased.


```
In [12]: catpath = '/content/drive/MyDrive/CVdata/grumpy-cat.png'
```

```
In [13]: # Applying the Kernel

# Reloading src and dst
src = cv2.imread(catpath)
dst = cv2.cvtColor(src, cv2.COLOR_BGR2GRAY)

# Using the filter with our Kernel
new_img2 = cv2.filter2D(src, -1, kernel, dst)

# Display
cv2.imshow(new_img2)
```

Answer the following with analysis:

- What does `filter2d()` function do that resulted to this image above?
- Provide a comparison between this new image and the previous image that we were able to generate.

The `filter2d()` function is used to apply a filter to an image. `filter2d()` works by convolving the filter kernel with the image. This means that the kernel gets through the image and for each pixel, the kernel is multiplied by the corresponding pixel in the image, which will be summed. The sum we got will be assigned to the current pixel output of an image.

Based on the result, the `new_img2` is more sharpened and got a little blurring compared to the `src` image. We can also see that the output of the new images has a presence of noise, thinner edges and a slightly blurred look.

Based on this simple example, let's add two classes. One class, `VConvolutionFilter`, will represent a convolution filter in general. A subclass, `SharpenFilter`, will represent our sharpening filter specifically.

```
In [14]: class VConvolutionFilter(object):
        """a filter that applies a convolution to V (or all of BGR)."""

        def __init__(self, kernel):
            self._kernel = kernel

        def apply(self, src, dst):
            """ Apply the filter with a BGR or gray source/destination """
            cv2.filter2D(src, -1, self._kernel, dst)
```

```
In [15]: class SharpenFilter(VConvolutionFilter):
        """a sharpen filter with a 1-pixel radius."""

        def __init__(self):
```



```

kernel = np.array([[ -1, -1, -1],
                   [ -1,  9, -1],
                   [ -1, -1, -1]])
VConvolutionFilter.__init__(self, kernel)

```

Run the classes above, create objects and aim to show the output of the two classes. Afterwards, make sure to make an analysis of the output.

```

In [26]: # Reloading src and dst
src = cv2.imread(catpath)
dst = cv2.cvtColor(src, cv2.COLOR_BGR2GRAY)

```

```

In [31]: import matplotlib.pyplot as plt

# Apply the sharpened image
sharpen_filter = SharpenFilter()
dst_sharpened_color = np.zeros_like(src)
sharpen_filter.apply(src, dst_sharpened_color)

# Apply grayscale
src_gray = cv2.cvtColor(src, cv2.COLOR_BGR2GRAY) #convert the src to grayscale
dst_sharpened = np.zeros_like(src_gray)
sharpen_filter.apply(src_gray, dst_sharpened)

# Display the original and sharpened colored images
fig, axes = plt.subplots(1, 2, figsize=(10, 5))
axes[0].imshow(cv2.cvtColor(src, cv2.COLOR_BGR2RGB))
axes[0].set_title("Original Colored Image")
axes[0].axis("off")

axes[1].imshow(cv2.cvtColor(dst_sharpened_color, cv2.COLOR_BGR2RGB))
axes[1].set_title("Sharpened Colored Image")
axes[1].axis("off")


# Display the original and sharpened grayscale images
fig, axes = plt.subplots(1, 2, figsize=(10, 5))
axes[0].imshow(src_gray, cmap="gray")
axes[0].set_title("Original Grayscale")
axes[0].axis("off")

axes[1].imshow(dst_sharpened, cmap="gray")
axes[1].set_title("Sharpened Grayscale")
axes[1].axis("off")


plt.show()

```

Original Colored Image



Sharpened Colored Image



Original Grayscale



Sharpened Grayscale



Analysis

The original image is more natural than the sharpened image. The sharpened image enhances the edges of our images and produces noise.

Note that the weights sum up to 1. This should be the case whenever we want to leave the image's overall brightness unchanged. If we modify a sharpening kernel slightly so that its weights sum up to 0 instead, we have an edge detection kernel that turns edges white and non-edges black.

```
In [34]: class FindEdgesFilter(VConvolutionFilter):
  """An edge-finding filter with a 1-pixel radius."""

  def __init__(self):
    kernel = np.array([[ -1, -1, -1],
                       [-1, 8, -1],
                       [-1, -1, -1]])
    VConvolutionFilter.__init__(self, kernel)
```

Run this class and demonstrate the output. Provide an analysis

```
In [42]: findedges_filter = FindEdgesFilter()
dst_findedge_color = np.zeros_like(src)
findedges_filter.apply(src, dst_findedge_color)

# Display the original and Edge Finder images
fig, axes = plt.subplots(1, 2, figsize=(10, 5))
axes[0].imshow(cv2.cvtColor(src, cv2.COLOR_BGR2RGB))
axes[0].set_title("Original Colored Image")
axes[0].axis("off")

axes[1].imshow(cv2.cvtColor(dst_findedge_color, cv2.COLOR_BGR2RGB))
axes[1].set_title("Edge Finder Image")
axes[1].axis("off")
```

```
Out[42]: (-0.5, 1329.5, 997.5, -0.5)
```

Original Colored Image



Edge Finder Image



Analysis

The output image shows where are the edges located within a 1 pixel radius. The white color of the output image indicate the values of edges in each pixel.

Next, let's make a blur filter. Generally, for a blur effect, the weights should sum up to 1 and should be positive throughout the neighborhood. For example, we can take a simple average of the neighborhood as follows

```
In [43]: class Blurfilter(VConvolutionFilter):
        """A blur filter with a 2-pixel radius"""

        def __init__(self):
            kernel = np.array([[0.04, 0.04, 0.04, 0.04, 0.04],
                               [0.04, 0.04, 0.04, 0.04, 0.04],
                               [0.04, 0.04, 0.04, 0.04, 0.04],
                               [0.04, 0.04, 0.04, 0.04, 0.04],
                               [0.04, 0.04, 0.04, 0.04, 0.04]])
            VConvolutionFilter.__init__(self, kernel)
```

Run this class and demonstrate the output. Provide an analysis

```
In [51]: blur_filter = Blurfilter()
        dst_blur_color = np.zeros_like(src)
        blur_filter.apply(src, dst_blur_color)

        # Display the original and blur filter images
        fig, axes = plt.subplots(1, 2, figsize=(10, 5))
        axes[0].imshow(cv2.cvtColor(src, cv2.COLOR_BGR2RGB))
        axes[0].set_title("Original Colored Image")
        axes[0].axis("off")


        axes[1].imshow(cv2.cvtColor(dst_blur_color, cv2.COLOR_BGR2RGB))
        axes[1].set_title("Blur Filter Image")
        axes[1].axis("off")
```

```
Out[51]: (-0.5, 1329.5, 997.5, -0.5)
```


Original Colored Image



Blur Filter Image



Analysis

Using the blur filter in our original image, we can see that our output is got really blurred than the original. As shown on the output, the edges of our image is less visible and make it more smooth.

Our sharpening, edge detection, and blur filters use kernels that are highly symmetric. Sometimes, though, kernels with less symmetry produce an interesting effect. Let's consider a kernel that blurs on one side (with positive weights) and sharpens on the other (with negative weights). It will produce a ridged or embossed effect.

```
In [53]: class EmbossFilter(VConvolutionFilter):
        """An emboss filter with a 1-pixel radius."""

        def __init__(self):
            kernel = np.array([[-2, -1, 0],
                               [-1, 1, 1],
                               [0, 1, 2]])
            VConvolutionFilter.__init__(self, kernel)

In [56]: emboss_filter = EmbossFilter()
        dst_emboss_color = np.zeros_like(src)
        emboss_filter.apply(src,dst_emboss_color)

        # Display the original and emboss filter images
        fig, axes = plt.subplots(1, 2, figsize=(10, 5))
        axes[0].imshow(cv2.cvtColor(src, cv2.COLOR_BGR2RGB))
        axes[0].set_title("Original Colored Image")
        axes[0].axis("off")

        axes[1].imshow(cv2.cvtColor(dst_emboss_color, cv2.COLOR_BGR2RGB))
        axes[1].set_title("Emboss Filter Image")
        axes[1].axis("off")
```

Out[56]: (-0.5, 1329.5, 997.5, -0.5)

Original Colored Image



Emboss Filter Image



Run this class and demonstrate the output. Provide an analysis

Analysis

In using the Emboss Filter in the `src` image we demonstrate how an emboss work. We can observed that the embossed image slightly enhance the edges as seen from its light and dark ares. This effect create a depth in the picture and loss its softness from the original image.

Edge Detection with Canny

OpenCV also offers a very handy function called Canny (after the algorithm's inventor, John F. Canny), which is very popular not only because of its effectiveness, but also the simplicity of its implementation in an OpenCV program, as it is a one-liner:

```
In [62]: # import cv2
# import numpy as np
# from google.colab.patches import cv2_imshow

img = cv2.imread('/content/drive/MyDrive/CVdata/paimon.png', 0)
cv2.imwrite("canny.jpg", cv2.Canny(img, 200, 300))
```

Out[62]: True

```
In [63]: img_canny = cv2.imread('/content/canny.jpg')
cv2_imshow(img_canny)
```




The Canny edge detection algorithm is quite complex but also interesting: it's a five-step process that denoises the image with a Gaussian filter, calculates gradients, applies non maximum suppression (NMS) on edges, a double threshold on all the detected edges to eliminate false positives, and, lastly, analyzes all the edges and their connection to each other to keep the real edges and discard the weak ones.

Try it on your own image, do you agree that it's an effective edge detection algorithm? Demonstrate your samples before making a conclusion

```
In [65]: # Using Canny Edge to simple Images
killua = cv2.imread('/content/drive/MyDrive/CVdata/killua.png', 0)
cv2.imwrite("SimpleCanny.jpg", cv2.Canny(killua, 200, 300))
```

Out[65]: True


```
In [66]: img_killua = cv2.imread('/content/SimpleCanny.jpg')  
cv2.imshow(img_killua)
```



```
In [93]: #using Complex images  
doom = cv2.imread('/content/drive/MyDrive/CVdata/drdoom.png', 0)  
cv2.imwrite("ComplexCanny.jpg", cv2.Canny(doom, 200, 220))
```

Out[93]: True

```
In [94]: img_doom = cv2.imread('/content/ComplexCanny.jpg')  
cv2.imshow(img_doom)
```

Analysis

Canny edge Detection effectively outlined some simple strong edges of the images. This indicates that the algorithm is suitable for sharp contrast and structure edges. Meanwhile, in a more complex image with some weak and strong edges, the image output loses some fine details due to its threshold as indicated in the more complex image. Adjusting these parameters would likely refine the results.

Contour Detection


Another vital task in computer vision is contour detection, not only because of the obvious aspect of detecting contours of subjects contained in an image or video frame, but because of the derivative operations connected with identifying contours.

These operations are, namely, computing bounding polygons, approximating shapes, and generally calculating regions of interest, which considerably simplify interaction with image data because a rectangular region with NumPy is easily defined with an array slice. We will be using this technique a lot when exploring the concept of object detection (including faces) and object tracking.

```
In [95]: # import numpy as np
# import cv2
# from google.colab.patches import cv2_imshow
```

```
img = np.zeros((200,200), dtype=np.uint8)
img[50:150, 50:150] = 255
```

```
In [96]: cv2.imshow(img)
```




```
In [97]: ret, thresh = cv2.threshold(img, 127, 255, 0)
print(ret)
print(thresh)
```

```
127.0
[[0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 ...
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]]
```

```
In [98]: contours, hierarchy = cv2.findContours(thresh, cv2.RETR_TREE, cv2.CHAIN_APPROX_SIMPLE)
color = cv2.cvtColor(img, cv2.COLOR_GRAY2BGR)
```

```
In [99]: img = cv2.drawContours(color, contours, -1, (0, 255, 0), 2)
cv2.imshow(color)
```



What is indicated by the green box in the middle? Why are there no other green lines? Provided an analysis

The green box in the middle indicates the boundary of the image or ROI (region of interest). This green box identifies the boundary or outline of our image. There are no other green lines because we select certain contours like hierarchy which only one image we have and it met the criteria. In the exercise, we only have one primary object to be detected and it was highlighted.

Contours - bounding box, minimum area rectangle, and minimum enclosing circle

Finding the contours of a square is a simple task; irregular, skewed, and rotated shapes bring the best out of the `cv2.findContours` utility function of OpenCV. Consider this sample image:



In a real-life application, we would be most interested in determining the bounding box of the subject, its minimum enclosing rectangle, and its circle. The `cv2.findContours` function in conjunction with a few other OpenCV utilities makes this very easy to accomplish:

```
In [199... # import cv2
# import numpy as np

img = cv2.pyrDown(cv2.imread("/content/drive/MyDrive/CVdata/hammer.png", cv2.IMREAD_UNCHANGED))

ret, thresh = cv2.threshold(cv2.cvtColor(img.copy(), cv2.COLOR_BGR2GRAY), 127, 255, cv2.THRESH_BINARY)

In [200... cv2_imshow(img)
```



```
In [201... contours, hier = cv2.findContours(thresh, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
```

```
In [202... for c in contours:
    # Find the bounding box coordinates
    x, y, w, h = cv2.boundingRect(c)
    cv2.rectangle(img, (x,y), (x+w, y+h), (0, 255, 0), 2)

    # Find minimum area
    rect = cv2.minAreaRect(c)

    # Calculate coordinates of the minimum area rectangle
    box = cv2.boxPoints(rect)

    # Normalize coordinates to integers
    box = np.int0(box)

    # Draw contours
    cv2.drawContours(img, [box], 0, (0, 0, 255), 3)

    # Calculate center and radius of minimum enclosing circle
    (x, y), radius = cv2.minEnclosingCircle(c)

    # Cast to integers
    center = (int(x), int(y))
    radius = int(radius)

    # Draw the circle
    img = cv2.circle(img, center, radius, (0, 255, 0), 2)
```

```
<ipython-input-202-935b8ec75673>:13: DeprecationWarning: `np.int0` is a deprecated alias for `np.intp`. (Deprecated NumPy 1.24)
    box = np.int0(box)
```

```
In [203... cv2.drawContours(img, contours, -1, (255, 0, 0), 1)
cv2.imshow(img)
```



Explain what has happened so far.

```
In [204... x, y, w, h = cv2.boundingRect(c)
```

This is a pretty straightforward conversion of contour information to the (x, y) coordinates, plus the height and width of the rectangle. Drawing this rectangle is an easy task and can be done using this code:

```
In [205... cv2.rectangle(img, (x,y), (x+w, y+h), (0, 255, 0), 2)
```



```
Out[205... array([[ 0,  0,  0, 255],
          [ 0,  0,  0, 255],
          [ 0,  0,  0, 255],
          ...,
          [ 0,  0,  0, 255],
          [ 0,  0,  0, 255],
          [ 0,  0,  0, 255]],

          [[ 0,  0,  0, 255],
          [ 0,  0,  0, 255],
          [ 0,  0,  0, 255],
          ...,
          [ 0,  0,  0, 255],
          [ 0,  0,  0, 255],
          [ 0,  0,  0, 255]],

          [[ 0,  0,  0, 255],
          [ 0,  0,  0, 255],
          [ 0,  0,  0, 255],
          ...,
          [ 0,  0,  0, 255],
          [ 0,  0,  0, 255],
          [ 0,  0,  0, 255]],

          ...,

          [[ 0,  0,  0, 255],
          [ 0,  0,  0, 255],
          [ 0,  0,  0, 255],
          ...,
          [ 0,  0,  0, 255],
          [ 0,  0,  0, 255],
          [ 0,  0,  0, 255]],

          [[ 0,  0,  0, 255],
          [ 0,  0,  0, 255],
          [ 0,  0,  0, 255],
          ...,
          [ 0,  0,  0, 255],
          [ 0,  0,  0, 255],
          [ 0,  0,  0, 255]]], dtype=uint8)
```

```
In [206... cv2.imshow(img)
```



Secondly, let's calculate the minimum area enclosing the subject:

```
In [207... rect = cv2.minAreaRect(c)
box = cv2.boxPoints(rect)
box = np.int0(box)
```

```
<ipython-input-207-ceb1eab3a452>:3: DeprecationWarning: `np.int0` is a deprecated alias for `np.intp`. (Deprecated NumPy 1.24)
    box = np.int0(box)
```

The mechanism used here is particularly interesting: OpenCV does not have a function to calculate the coordinates of the minimum rectangle vertexes directly from the contour information. Instead, we calculate the minimum rectangle area, and then calculate the vertexes of this rectangle.

Note that the calculated vertexes are floats, but pixels are accessed with integers (you can't access a "portion" of a pixel), so we need to operate this conversion. Next, we draw the box, which gives us the perfect opportunity to introduce the `cv2.drawContours` function:

```
In [208... cv2.drawContours(img, [box], 0, (0, 0, 255), 3)
```

```
Out[208... array([[ 0,  0,  0, 255],
        [ 0,  0,  0, 255],
        [ 0,  0,  0, 255],
        ...,
        [ 0,  0,  0, 255],
        [ 0,  0,  0, 255],
        [ 0,  0,  0, 255]],

       [[ 0,  0,  0, 255],
        [ 0,  0,  0, 255],
        [ 0,  0,  0, 255],
        ...,
        [ 0,  0,  0, 255],
        [ 0,  0,  0, 255],
        [ 0,  0,  0, 255]],

       [[ 0,  0,  0, 255],
        [ 0,  0,  0, 255],
        [ 0,  0,  0, 255],
        ...,
        [ 0,  0,  0, 255],
        [ 0,  0,  0, 255],
        [ 0,  0,  0, 255]],

       ...,

       [[ 0,  0,  0, 255],
        [ 0,  0,  0, 255],
        [ 0,  0,  0, 255],
        ...,
        [ 0,  0,  0, 255],
        [ 0,  0,  0, 255],
        [ 0,  0,  0, 255]],

       [[ 0,  0,  0, 255],
        [ 0,  0,  0, 255],
        [ 0,  0,  0, 255],
        ...,
        [ 0,  0,  0, 255],
        [ 0,  0,  0, 255],
        [ 0,  0,  0, 255]]], dtype=uint8)
```

```
In [209... cv2.imshow(img)
```



Firstly, this function—like all drawing functions—modifies the original image. Secondly, it takes an array of contours in its second parameter, so you can draw a number of contours in a single operation. Therefore, if you have a single set of points representing a contour polygon, you need to wrap these points into an array, exactly like we did with our box in the preceding example. The third parameter of this function specifies the index of the contours array that we want to draw: a value of -1 will draw all contours; otherwise, a contour at the specified index in the contours array (the second parameter) will be drawn.

Most drawing functions take the color of the drawing and its thickness as the last two parameters.

The last bounding contour we're going to examine is the minimum enclosing circle:

```
In [210... (x, y), radius = cv2.minEnclosingCircle(c)
center = (int(x), int(y))
radius = int(radius)
img = cv2.circle(img, center, radius, (0, 255, 0), 2)
```

The only peculiarity of the `cv2.minEnclosingCircle` function is that it returns a two-element tuple, of which the first element is a tuple itself, representing the coordinates of the circle's center, and the second element is the radius of this circle. After converting all these values to integers, drawing the circle is quite a trivial operation.

Show the final image output and provide an analysis. Does it look similar to the image shown below?



```
In [213... cv2.imshow(img)
```



```
In [212... plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
plt.axis("off")
```

```
Out[212... (-0.5, 129.5, 133.5, -0.5)
```




Analysis

The output we got in the activity is similar to the expected output shown above. The code provide a simple shaped contours for the hammer image.

4. Supplementary Activity

For this section of the activity, you must have your favorite fictional character's picture ready.

Perform/Answer the following:

- Run all classes above meant to filter an image to your favorite character.
- Use edge detection and contour detection on your fave character. Do they indicate the same?
- Modify your character's picture such that bounding boxes similar to what happens in the last procedure will be written on the image.
- Research on the benefits of using canny and contour detection. What happens to the image after edge detection? What happens when you apply contour straight after?

Filter Application to Dr. Doom

```
In [214... src = cv2.imread('/content/drive/MyDrive/CVdata/drdoom.png')  
dst = cv2.cvtColor(src, cv2.COLOR_BGR2GRAY)
```

```
In [215... cv2.imshow(src)
```



In [216... `cv2_imshow(dst)`



Stroke Edges Sample Output

```
In [228... new_doom_img = strokeEdges(src,dst)
new_blur_doom_img = strokeEdges(src,dst,blurKsize= 2)
new_edge_doom_img = strokeEdges(src,dst,edgeKsize= 15)

fig, axes = plt.subplots(1, 4, figsize=(10, 5))
axes[0].imshow(cv2.cvtColor(src, cv2.COLOR_BGR2RGB))
axes[0].set_title("Original Colored Image")
axes[0].axis("off")

axes[1].imshow(cv2.cvtColor(new_doom_img, cv2.COLOR_BGR2RGB))
axes[1].set_title("Stroke Edges Image")
axes[1].axis("off")

axes[2].imshow(cv2.cvtColor(new_blur_doom_img, cv2.COLOR_BGR2RGB))
axes[2].set_title("SEI(blurKsize = 2)")
axes[2].axis("off")

axes[3].imshow(cv2.cvtColor(new_edge_doom_img, cv2.COLOR_BGR2RGB))
axes[3].set_title("SEI(edgeKsize = 15)")
axes[3].axis("off")
```

```
Out[228... (-0.5, 443.5, 691.5, -0.5)
```


Original Colored Image




Stroke Edges Image



SEI(blurKsize = 2)



SEI(edgeKsize = 15)



filter2d() Sample Output

```
In [231...] new_doom_img2 = cv2.filter2D(src, -1, kernel, dst)
```

```
In [232...] cv2.imshow(new_doom_img2)
```



Shapen Filter Output

```
In [233... sharpen_filter = SharpenFilter()

sharpened_img = src.copy()
sharpen_filter.apply(src,sharpened_img)

# Display the original and sharpened colored images
fig, axes = plt.subplots(1, 2, figsize=(10, 5))
axes[0].imshow(cv2.cvtColor(src, cv2.COLOR_BGR2RGB))
axes[0].set_title("Original Colored Image")
axes[0].axis("off")

axes[1].imshow(cv2.cvtColor(sharpened_img, cv2.COLOR_BGR2RGB))
axes[1].set_title("Sharpened Colored Image")
axes[1].axis("off")
```

Out[233... (-0.5, 443.5, 691.5, -0.5)

Original Colored Image



Sharpened Colored Image



Find Edge Filter Output

```
In [236... findedges_filter = FindEdgesFilter()

findedges_img = src.copy()
findedges_filter.apply(src,findedges_img)

# Display the original and edge filter images
fig, axes = plt.subplots(1, 2, figsize=(10, 5))
axes[0].imshow(cv2.cvtColor(src, cv2.COLOR_BGR2RGB))
axes[0].set_title("Original Colored Image")
axes[0].axis("off")

axes[1].imshow(cv2.cvtColor(findedges_img, cv2.COLOR_BGR2RGB))
axes[1].set_title("Find Edge Image")
axes[1].axis("off")
```

Out[236... (-0.5, 443.5, 691.5, -0.5)

Original Colored Image



Find Edge Image



Blur Filter Ouput

```
In [237...] blur_filter = Blurfilter()

blur_img = src.copy()
blur_filter.apply(src, blur_img)

# Display the original and blur filter images
fig, axes = plt.subplots(1, 2, figsize=(10, 5))
axes[0].imshow(cv2.cvtColor(src, cv2.COLOR_BGR2RGB))
axes[0].set_title("Original Image")
axes[0].axis("off")


axes[1].imshow(cv2.cvtColor(blur_img, cv2.COLOR_BGR2RGB))
axes[1].set_title("Blur Filter Image")
axes[1].axis("off")
```

```
Out[237...] (-0.5, 443.5, 691.5, -0.5)
```


Original Image



Blur Filter Image



Emboss Filter Image

```
In [238... emboss_filter = EmbossFilter()

emboss_img = src.copy()
emboss_filter.apply(src, emboss_img)

# Display the original and emboss filter images
fig, axes = plt.subplots(1, 2, figsize=(10, 5))
axes[0].imshow(cv2.cvtColor(src, cv2.COLOR_BGR2RGB))
axes[0].set_title("Original Colored Image")
axes[0].axis("off")

axes[1].imshow(cv2.cvtColor(emboss_img, cv2.COLOR_BGR2RGB))
axes[1].set_title("Emboss Filter Image")
axes[1].axis("off")
```

Out[238... (-0.5, 443.5, 691.5, -0.5)

Original Colored Image



Emboss Filter Image




Contour and Edge Application to Dr.Doom

Edge

```
In [270...] doom_img = cv2.imread('/content/drive/MyDrive/CVdata/drdoom.png')  
cv2.imwrite('DrDoomCanny.jpg', cv2.Canny(doom_img, 150, 200))
```

Out[270...] True

```
In [271...] img_doom = cv2.imread('/content/DrDoomCanny.jpg')  
cv2.imshow(img_doom)
```




Contour


```
In [282... img = cv2.imread('/content/DrDoomCanny.jpg', cv2.IMREAD_GRAYSCALE)
cv2.imshow(img)

# Apply thresholding to the grayscale image 'img'
ret, thresh = cv2.threshold(img, 127, 255, 0)

# Now 'thresh' will be a single-channel image suitable for findContours
contours, hierarchy = cv2.findContours(thresh, cv2.RETR_TREE, cv2.CHAIN_APPROX_SIMPLE)
color = cv2.cvtColor(img, cv2.COLOR_GRAY2BGR)
```

```
In [283... img = cv2.drawContours(color, contours, -1, (0, 255, 0), 2)
cv2.imshow(color)
```




```
In [286... # Display the Edge and Contour images
fig, axes = plt.subplots(1, 2, figsize=(10, 5))
axes[0].imshow(cv2.cvtColor(img_doom, cv2.COLOR_BGR2RGB))
axes[0].set_title("Edge Image")
axes[0].axis("off")

axes[1].imshow(cv2.cvtColor(color, cv2.COLOR_BGR2RGB))
axes[1].set_title("Contour Image")
axes[1].axis("off")
```

```
Out[286... (-0.5, 443.5, 691.5, -0.5)
```

Edge Image



Contour Image



Contour with Bounding Box Application

```
In [290...] img = cv2.pyrDown(cv2.imread('/content/drive/MyDrive/CVdata/drdoom.png'), cv2.IMREAD_GRAYSCALE)
ret, thresh = cv2.threshold(cv2.cvtColor(img.copy(), cv2.COLOR_BGR2GRAY), 127, 255, cv2.THRESH_BINARY)
```

```
In [291...] cv2.imshow('img')
```



```
In [292...] contours, hier = cv2.findContours(thresh, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
```

```
In [295...] for c in contours:
    # Find the bounding box coordinates
    x, y, w, h = cv2.boundingRect(c)
    cv2.rectangle(img, (x,y), (x+w, y+h), (0, 255, 0), 2)

    # Find minimum area
    rect = cv2.minAreaRect(c)

    # Calculate coordinates of the minimum area rectangle
    box = cv2.boxPoints(rect)
```

```

# Normalize coordinates to integers
box = np.int0(box)

# Draw contours
cv2.drawContours(img, [box], 0, (0, 0, 255), 3)

# Calculate center and radius of minimum enclosing circle
(x, y), radius = cv2.minEnclosingCircle(c)

# Cast to integers
center = (int(x), int(y))
radius = int(radius)

# Draw the circle
img = cv2.circle(img, center, radius, (0, 255, 0), 2)

```

```

<ipython-input-295-7125e5b9682f>:13: DeprecationWarning: `np.int0` is a deprecated alias for `np.intp`. (Deprecated NumPy 1.24)
box = np.int0(box)

```

```

In [296... cv2.drawContours(img, contours, -1, (255, 0, 0), 1)
cv2.imshow(img)

```



Benefits of using canny and contour detection.

Canny Edge detection is a technique on image processing which detect the edges in images. It recognize both strong and weak lines of the image which is good on outlining shapes and details. It also reduce variation of brightness and color, preventing false edges during the process. Meanwhile, Contour detection is use to identify or recognize object. It is useful for maching shapes within an image which are essentially usefu for object detection.

After edge detection, the image transformed into binary where edges are highlighted. During this process, it will help use indentify the boundaries of an object and edges represent thin lines that outline the shape of an image.

Applying contour detection straight after edge detection will identify and outlined the contours of our object. It will trace the edge detected and form a continuous line which represent the boundary of our ROI resulting to more detailed object representation within an image.

5. Summary, Conclusions and Lessons Learned

In this notebook we perform various task on edge and contour detection as well as image processing manipulation such as filter, blur and edge. We also mänge to use the apply what we learn in the exercises in our supplementary activity. I also compare the original image from the filter we used. The challenging part of this activity is changing the parameters of

contours and edges to detect the object since my image is complex which is hard for the program to identify its edge lines.

Proprietary Clause

Property of the Technological Institute of the Philippines (T.I.P.). No part of the materials made and uploaded in this learning management system by T.I.P. may be copied, photographed, printed, reproduced, shared, transmitted, translated, or reduced to any electronic medium or machine-readable form, in whole or in part, without the prior consent of T.I.P.